

O'REILLY®

Hadoop with Python



Zachary Radtka
& Donald Miner

Additional Resources

4 Easy Ways to Learn More and Stay Current

Programming Newsletter

Get programming related news and content delivered weekly to your inbox.

oreilly.com/programming/newsletter

Free Webcast Series

Learn about popular programming topics from experts live, online.

webcasts.oreilly.com

O'Reilly Radar

Read more insight and analysis about emerging technologies.

radar.oreilly.com

Conferences

Immerse yourself in learning at an upcoming O'Reilly conference.

conferences.oreilly.com

Hadoop with Python

Zachary Radtka & Donald Miner

Beijing • Boston • Farnham • Sebastopol • Tokyo

O'REILLY®

Hadoop with Python

by Zachary Radtka and Donald Miner

Copyright © 2016 O'Reilly Media, Inc.. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://safaribooksonline.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Editor: Meghan Blanchette

Production Editor: Kristen Brown

Copyeditor: Sonia Saruba

Interior Designer: David Futato

Cover Designer: Karen Montgomery

Illustrator: Rebecca Demarest

October 2015: First Edition

Revision History for the First Edition

2015-10-19 First Release

See <http://oreilly.com/catalog/errata.csp?isbn=9781491942277> for release details.

While the publisher and the authors have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the authors disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-491-94227-7

[LSI]

Table of Contents

Source Code.....	vii
1. Hadoop Distributed File System (HDFS).....	1
Overview of HDFS	2
Interacting with HDFS	3
Snakebite	7
Chapter Summary	13
2. MapReduce with Python.....	15
Data Flow	15
Hadoop Streaming	18
mrjob	22
Chapter Summary	26
3. Pig and Python.....	27
WordCount in Pig	28
Running Pig	29
Pig Latin	31
Extending Pig with Python	35
Chapter Summary	40
4. Spark with Python.....	41
WordCount in PySpark	41
PySpark	43
Resilient Distributed Datasets (RDDs)	44
Text Search with PySpark	50

Chapter Summary	52
5. Workflow Management with Python.....	53
Installation	53
Workflows	54
An Example Workflow	55
Hadoop Workflows	58
Chapter Summary	62

Source Code

All of the source code in this book is on GitHub. To copy the source code locally, use the following `git clone` command:

```
$ git clone https://github.com/MinerKasch/HadoopWithPython
```


Hadoop Distributed File System (HDFS)

The Hadoop Distributed File System (HDFS) is a Java-based distributed, scalable, and portable filesystem designed to span large clusters of commodity servers. The design of HDFS is based on GFS, the Google File System, which is described in a [paper](#) published by Google. Like many other distributed filesystems, HDFS holds a large amount of data and provides transparent access to many clients distributed across a network. Where HDFS excels is in its ability to store very large files in a reliable and scalable manner.

HDFS is designed to store a lot of information, typically petabytes (for very large files), gigabytes, and terabytes. This is accomplished by using a block-structured filesystem. Individual files are split into fixed-size blocks that are stored on machines across the cluster. Files made of several blocks generally do not have all of their blocks stored on a single machine.

HDFS ensures reliability by replicating blocks and distributing the replicas across the cluster. The default replication factor is three, meaning that each block exists three times on the cluster. Block-level replication enables data availability even when machines fail.

This chapter begins by introducing the core concepts of HDFS and explains how to interact with the filesystem using the native built-in commands. After a few examples, a Python client library is introduced that enables HDFS to be accessed programmatically from within Python applications.

Overview of HDFS

The architectural design of HDFS is composed of two processes: a process known as the NameNode holds the metadata for the filesystem, and one or more DataNode processes store the blocks that make up the files. The NameNode and DataNode processes can run on a single machine, but HDFS clusters commonly consist of a dedicated server running the NameNode process and possibly thousands of machines running the DataNode process.

The NameNode is the most important machine in HDFS. It stores metadata for the entire filesystem: filenames, file permissions, and the location of each block of each file. To allow fast access to this information, the NameNode stores the entire metadata structure in memory. The NameNode also tracks the replication factor of blocks, ensuring that machine failures do not result in data loss. Because the NameNode is a single point of failure, a secondary NameNode can be used to generate snapshots of the primary NameNode's memory structures, thereby reducing the risk of data loss if the NameNode fails.

The machines that store the blocks within HDFS are referred to as DataNodes. DataNodes are typically commodity machines with large storage capacities. Unlike the NameNode, HDFS will continue to operate normally if a DataNode fails. When a DataNode fails, the NameNode will replicate the lost blocks to ensure each block meets the minimum replication factor.

The example in **Figure 1-1** illustrates the mapping of files to blocks in the NameNode, and the storage of blocks and their replicas within the DataNodes.

The following section describes how to interact with HDFS using the built-in commands.

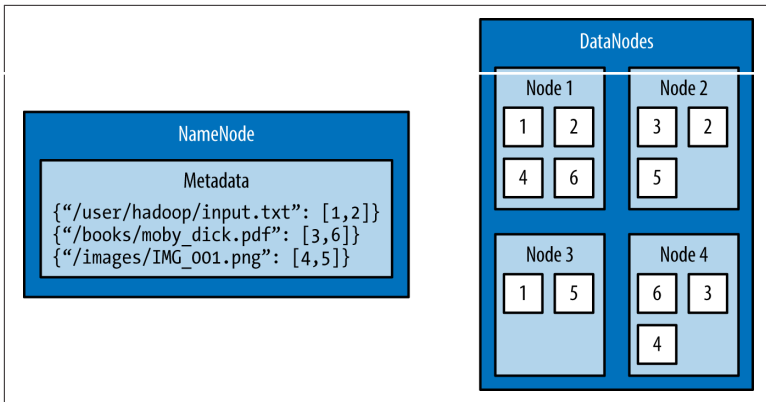


Figure 1-1. An HDFS cluster with a replication factor of two; the NameNode contains the mapping of files to blocks, and the DataNodes store the blocks and their replicas

Interacting with HDFS

Interacting with HDFS is primarily performed from the command line using the script named `hdfs`. The `hdfs` script has the following usage:

```
$ hdfs COMMAND [-option <arg>]
```

The `COMMAND` argument instructs which functionality of HDFS will be used. The `-option` argument is the name of a specific option for the specified command, and `<arg>` is one or more arguments that are specified for this option.

Common File Operations

To perform basic file manipulation operations on HDFS, use the `dfs` command with the `hdfs` script. The `dfs` command supports many of the same file operations found in the Linux shell.

It is important to note that the `hdfs` command runs with the permissions of the system user running the command. The following examples are run from a user named “`hduser`.”

List Directory Contents

To list the contents of a directory in HDFS, use the `-ls` command:

```
$ hdfs dfs -ls
$
```

Running the `-ls` command on a new cluster will not return any results. This is because the `-ls` command, without any arguments, will attempt to display the contents of the user's home directory on HDFS. This is not the same home directory on the host machine (e.g., `/home/$USER`), but is a directory within HDFS.

Providing `-ls` with the forward slash (`/`) as an argument displays the contents of the root of HDFS:

```
$ hdfs dfs -ls /
Found 2 items
drwxr-xr-x - hadoop supergroup    0 2015-09-20 14:36 /hadoop
drwx----- - hadoop supergroup    0 2015-09-20 14:36 /tmp
```

The output provided by the `hdfs dfs` command is similar to the output on a Unix filesystem. By default, `-ls` displays the file and folder permissions, owners, and groups. The two folders displayed in this example are automatically created when HDFS is formatted. The `hadoop` user is the name of the user under which the Hadoop daemons were started (e.g., `NameNode` and `DataNode`), and the `supergroup` is the name of the group of superusers in HDFS (e.g., `hadoop`).

Creating a Directory

Home directories within HDFS are stored in `/user/$HOME`. From the previous example with `-ls`, it can be seen that the `/user` directory does not currently exist. To create the `/user` directory within HDFS, use the `-mkdir` command:

```
$ hdfs dfs -mkdir /user
```

To make a home directory for the current user, `hduser`, use the `-mkdir` command again:

```
$ hdfs dfs -mkdir /user/hduser
```

Use the `-ls` command to verify that the previous directories were created:

```
$ hdfs dfs -ls -R /user
drwxr-xr-x - hduser supergroup    0 2015-09-22 18:01 /user/
hduser
```

Copy Data onto HDFS

After a directory has been created for the current user, data can be uploaded to the user's HDFS home directory with the `-put` command:

```
$ hdfs dfs -put /home/hduser/input.txt /user/hduser
```

This command copies the file `/home/hduser/input.txt` from the local filesystem to `/user/hduser/input.txt` on HDFS.

Use the `-ls` command to verify that `input.txt` was moved to HDFS:

```
$ hdfs dfs -ls
Found 1 items
-rw-r--r--    1 hduser supergroup          52 2015-09-20 13:20
input.txt
```

Retrieving Data from HDFS

Multiple commands allow data to be retrieved from HDFS. To simply view the contents of a file, use the `-cat` command. `-cat` reads a file on HDFS and displays its contents to stdout. The following command uses `-cat` to display the contents of `/user/hduser/input.txt`:

```
$ hdfs dfs -cat input.txt
jack be nimble
jack be quick
jack jumped over the candlestick
```

Data can also be copied from HDFS to the local filesystem using the `-get` command. The `-get` command is the opposite of the `-put` command:

```
$ hdfs dfs -get input.txt /home/hduser
```

This command copies `input.txt` from `/user/hduser` on HDFS to `/home/hduser` on the local filesystem.

HDFS Command Reference

The commands demonstrated in this section are the basic file operations needed to begin using HDFS. Below is a full listing of file manipulation commands possible with `hdfs dfs`. This listing can also be displayed from the command line by specifying `hdfs dfs` without any arguments. To get help with a specific option, use either `hdfs dfs -usage <option>` or `hdfs dfs -help <option>`.

```

Usage: hadoop fs [generic options]
    [-appendToFile <localsrc> ... <dst>]
    [-cat [-ignoreCrc] <src> ...]
    [-checksum <src> ...]
    [-chgrp [-R] GROUP PATH...]
    [-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
    [-chown [-R] [OWNER][:[GROUP]] PATH...]
    [-copyFromLocal [-f] [-p] [-l] <localsrc> ... <dst>]
        [-copyToLocal [-p] [-ignoreCrc] [-crc] <src> ...
<localdst>]
    [-count [-q] [-h] <path> ...]
    [-cp [-f] [-p | -p[topax]] <src> ... <dst>]
    [-createSnapshot <snapshotDir> [<snapshotName>]]
    [-deleteSnapshot <snapshotDir> <snapshotName>]
    [-df [-h] [<path> ...]]
    [-du [-s] [-h] <path> ...]
    [-expunge]
    [-find <path> ... <expression> ...]
    [-get [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
    [-getfacl [-R] <path>]
    [-getfattr [-R] {-n name | -d} [-e en] <path>]
    [-getmerge [-nl] <src> <localdst>]
    [-help [cmd ...]]
    [-ls [-d] [-h] [-R] [<path> ...]]
    [-mkdir [-p] <path> ...]
    [-moveFromLocal <localsrc> ... <dst>]
    [-moveToLocal <src> <localdst>]
    [-mv <src> ... <dst>]
    [-put [-f] [-p] [-l] <localsrc> ... <dst>]
    [-renameSnapshot <snapshotDir> <oldName> <newName>]
    [-rm [-f] [-r|-R] [-skipTrash] <src> ...]
    [-rmdir [--ignore-fail-on-non-empty] <dir> ...]
    [-setfacl [-R] [{-b|-k} {-m|-x <acl_spec>} <path>}][--set
<acl_spec> <path>]]
    [-setfattr {-n name [-v value] | -x name} <path>]
    [-setrep [-R] [-w] <rep> <path> ...]
    [-stat [format] <path> ...]
    [-tail [-f] <file>]
    [-test [-defsz] <path>]
    [-text [-ignoreCrc] <src> ...]
    [-touchz <path> ...]
    [-truncate [-w] <length> <path> ...]
    [-usage [cmd ...]]

```

Generic options supported are

```

-conf <configuration file>      specify an application configura-
tion file
-D <property=value>             use value for given property
-fs <local|namenode:port>       specify a namenode
-jt <local|resourcemanager:port> specify a ResourceManager
-files <comma separated list of files> specify comma separa-

```

```
ted files to be copied to the map reduce cluster
-libjars <comma separated list of jars>      specify comma sepa-
rated jar files to include in the classpath.
-archives <comma separated list of archives>  specify comma
separated archives to be unarchived on the compute machines.
```

```
The general command line syntax is
bin/hadoop command [genericOptions] [commandOptions]
```

The next section introduces a Python library that allows HDFS to be accessed from within Python applications.

Snakebite

Snakebite is a Python package, created by Spotify, that provides a Python client library, allowing HDFS to be accessed programmatically from Python applications. The client library uses protobuf messages to communicate directly with the NameNode. The Snakebite package also includes a command-line interface for HDFS that is based on the client library.

This section describes how to install and configure the Snakebite package. Snakebite's client library is explained in detail with multiple examples, and Snakebite's built-in CLI is introduced as a Python alternative to the `hdfs dfs` command.

Installation

Snakebite requires Python 2 and python-protobuf 2.4.1 or higher. Python 3 is currently not supported.

Snakebite is distributed through PyPI and can be installed using `pip`:

```
$ pip install snakebite
```

Client Library

The client library is written in Python, uses protobuf messages, and implements the Hadoop RPC protocol for talking to the NameNode. This enables Python applications to communicate directly with HDFS and not have to make a system call to `hdfs dfs`.

List Directory Contents

Example 1-1 uses the Snakebite client library to list the contents of the root directory in HDFS.

Example 1-1. *python/HDFS/list_directory.py*

```
from snakebite.client import Client

client = Client('localhost', 9000)
for x in client.ls(['/']):
    print x
```

The most important line of this program, and every program that uses the client library, is the line that creates a client connection to the HDFS NameNode:

```
client = Client('localhost', 9000)
```

The `Client()` method accepts the following parameters:

`host` (*string*)

Hostname or IP address of the NameNode

`port` (*int*)

RPC port of the NameNode

`hadoop_version` (*int*)

The Hadoop protocol version to be used (default: 9)

`use_trash` (*boolean*)

Use trash when removing files

`effective_use` (*string*)

Effective user for the HDFS operations (default: None or current user)

The `host` and `port` parameters are required and their values are dependent upon the HDFS configuration. The values for these parameters can be found in the *hadoop/conf/core-site.xml* configuration file under the property `fs.defaultFS`:

```
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:9000</value>
</property>
```

For the examples in this section, the values used for `host` and `port` are `localhost` and `9000`, respectively.

After the client connection is created, the HDFS filesystem can be accessed. The remainder of the previous application used the `ls` command to list the contents of the root directory in HDFS:


```
for x in client.ls(['/']):
    print x
```

It is important to note that many of methods in Snakebite return generators. Therefore they must be consumed to execute. The `ls` method takes a list of paths and returns a list of maps that contain the file information.

Executing the *list_directory.py* application yields the following results:

```
$ python list_directory.py
{'group': u'supergroup', 'permission': 448, 'file_type': 'd',
 'access_time': 0L, 'block_replication': 0, 'modification_time': 1442752574936L, 'length': 0L, 'blocksize': 0L,
 'owner': u'hduuser', 'path': '/tmp'}
{'group': u'supergroup', 'permission': 493, 'file_type': 'd',
 'access_time': 0L, 'block_replication': 0, 'modification_time': 1442742056276L, 'length': 0L, 'blocksize': 0L,
 'owner': u'hduuser', 'path': '/user'}
```

Create a Directory

Use the `mkdir()` method to create directories on HDFS. **Example 1-2** creates the directories */foo/bar* and */input* on HDFS.

Example 1-2. python/HDFS/mkdir.py

```
from snakebite.client import Client

client = Client('localhost', 9000)
for p in client.mkdir(['/foo/bar', '/input'], create_parent=True):
    print p
```

Executing the *mkdir.py* application produces the following results:

```
$ python mkdir.py
{'path': '/foo/bar', 'result': True}
{'path': '/input', 'result': True}
```

The `mkdir()` method takes a list of paths and creates the specified paths in HDFS. This example used the `create_parent` parameter to ensure that parent directories were created if they did not already exist. Setting `create_parent` to `True` is analogous to the `mkdir -p` Unix command.

Deleting Files and Directories

Deleting files and directories from HDFS can be accomplished with the `delete()` method. **Example 1-3** recursively deletes the `/foo` and `/bar` directories, created in the previous example.

Example 1-3. python/HDFS/delete.py

```
from snakebite.client import Client

client = Client('localhost', 9000)
for p in client.delete(['/foo', '/input'], recurse=True):
    print p
```

Executing the `delete.py` application produces the following results:

```
$ python delete.py
{'path': '/foo', 'result': True}
{'path': '/input', 'result': True}
```

Performing a recursive delete will delete any subdirectories and files that a directory contains. If a specified path cannot be found, the delete method throws a `FileNotFoundException`. If `recurse` is not specified and a subdirectory or file exists, `DirectoryException` is thrown.

The `recurse` parameter is equivalent to `rm -rf` and should be used with care.

Retrieving Data from HDFS

Like the `hdfs dfs` command, the client library contains multiple methods that allow data to be retrieved from HDFS. To copy files from HDFS to the local filesystem, use the `copyToLocal()` method. **Example 1-4** copies the file `/input/input.txt` from HDFS and places it under the `/tmp` directory on the local filesystem.

Example 1-4. python/HDFS/copy_to_local.py

```
from snakebite.client import Client

client = Client('localhost', 9000)
for f in client.copyToLocal(['/input/input.txt'], '/tmp'):
    print f
```

Executing the `copy_to_local.py` application produces the following result:

```
$ python copy_to_local.py
{'path': '/tmp/input.txt', 'source_path': '/input/input.txt',
 'result': True, 'error': ''}
```

To simply read the contents of a file that resides on HDFS, the `text()` method can be used. [Example 1-5](#) displays the content of `/input/input.txt`.

Example 1-5. `python/HDFS/text.py`

```
from snakebite.client import Client

client = Client('localhost', 9000)
for l in client.text(['/input/input.txt']):
    print l
```

Executing the `text.py` application produces the following results:

```
$ python text.py
jack be nimble
jack be quick
jack jumped over the candlestick
```

The `text()` method will automatically uncompress and display gzip and bzip2 files.

CLI Client

The CLI client included with Snakebite is a Python command-line HDFS client based on the client library. To execute the Snakebite CLI, the hostname or IP address of the NameNode and RPC port of the NameNode must be specified. While there are many ways to specify these values, the easiest is to create a `~/.snakebiterc` configuration file. [Example 1-6](#) contains a sample config with the NameNode hostname of `localhost` and RPC port of `9000`.

Example 1-6. `~/.snakebiterc`

```
{
    "config_version": 2,
    "skiptrash": true,
    "namenodes": [
        {"host": "localhost", "port": 9000, "version": 9},
    ]
}
```

The values for `host` and `port` can be found in the `hadoop/conf/core-site.xml` configuration file under the property `fs.defaultFS`.

For more information on configuring the CLI, see the [Snakebite CLI documentation online](#).

Usage

To use the Snakebite CLI client from the command line, simply use the command `snakebite`. Use the `ls` option to display the contents of a directory:

```
$ snakebite ls /
Found 2 items
drwx----- - hadoop    supergroup    0 2015-09-20 14:36 /tmp
drwxr-xr-x - hadoop    supergroup    0 2015-09-20 11:40 /user
```

Like the `hdfs dfs` command, the CLI client supports many familiar file manipulation commands (e.g., `ls`, `mkdir`, `df`, `du`, etc.).

The major difference between `snakebite` and `hdfs dfs` is that `snakebite` is a pure Python client and does not need to load any Java libraries to communicate with HDFS. This results in quicker interactions with HDFS from the command line.

CLI Command Reference

The following is a full listing of file manipulation commands possible with the `snakebite` CLI client. This listing can be displayed from the command line by specifying `snakebite` without any arguments. To view help with a specific command, use `snakebite [cmd] --help`, where `cmd` is a valid `snakebite` command.

```
snakebite [general options] cmd [arguments]
general options:
  -D --debug           Show debug information
  -V --version         Hadoop protocol version (default:9)
  -h --help           show help
  -j --json           JSON output
  -n --namenode       namenode host
  -p --port           namenode RPC port (default: 8020)
  -v --ver            Display snakebite version

commands:
  cat [paths]         copy source paths to stdout
  chgrp <grp> [paths] change group
  chmod <mode> [paths] change file mode (octal)
  chown <owner:grp> [paths] change owner
  copyToLocal [paths] dst copy paths to local
```

<code>count [paths]</code>	file system destination
<code>df</code>	display stats for paths
<code>du [paths]</code>	display fs stats
<code>get file dst</code>	display disk usage statistics
	copy files to local
	file system destination
<code>getmerge dir dst</code>	concatenates files in source dir
	into destination local file
<code>ls [paths]</code>	list a path
<code>mkdir [paths]</code>	create directories
<code>mkdirp [paths]</code>	create directories and their
	parents
<code>mv [paths] dst</code>	move paths to destination
<code>rm [paths]</code>	remove paths
<code>rmdir [dirs]</code>	delete a directory
<code>serverdefaults</code>	show server information
<code>setrep <rep> [paths]</code>	set replication factor
<code>stat [paths]</code>	stat information
<code>tail path</code>	display last kilobyte of the
	file to stdout
<code>test path</code>	test a path
<code>text path [paths]</code>	output file in text format
<code>touchz [paths]</code>	creates a file of zero length
<code>usage <cmd></code>	show cmd usage

to see command-specific options use: `snakebite [cmd] --help`

Chapter Summary

This chapter introduced and described the core concepts of HDFS. It explained how to interact with the filesystem using the built-in `hdfs dfs` command. It also introduced the Python library, Snakebite. Snakebite's client library was explained in detail with multiple examples. The snakebite CLI was also introduced as a Python alternative to the `hdfs dfs` command.

MapReduce with Python

MapReduce is a programming model that enables large volumes of data to be processed and generated by dividing work into independent tasks and executing the tasks in parallel across a cluster of machines. The MapReduce programming style was inspired by the functional programming constructs `map` and `reduce`, which are commonly used to process lists of data. At a high level, every MapReduce program transforms a list of input data elements into a list of output data elements twice, once in the `map` phase and once in the `reduce` phase.

This chapter begins by introducing the MapReduce programming model and describing how data flows through the different phases of the model. Examples then show how MapReduce jobs can be written in Python.

Data Flow

The MapReduce framework is composed of three major phases: `map`, `shuffle` and `sort`, and `reduce`. This section describes each phase in detail.

Map

The first phase of a MapReduce application is the `map` phase. Within the `map` phase, a function (called the `mapper`) processes a series of key-value pairs. The `mapper` sequentially processes each

key-value pair individually, producing zero or more output key-value pairs (**Figure 2-1**).

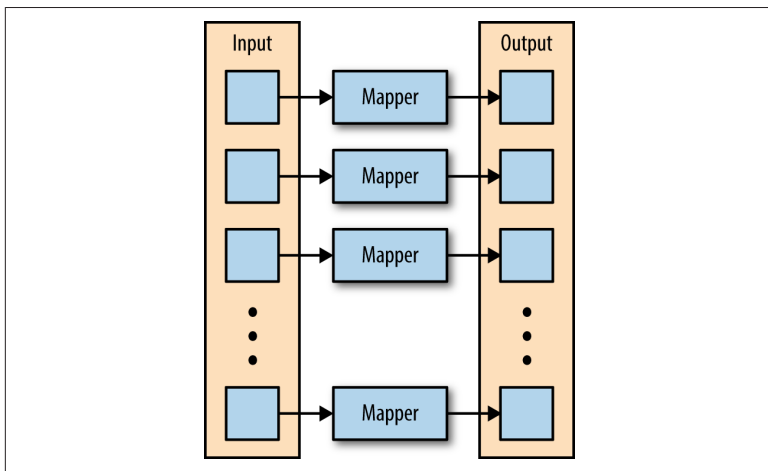


Figure 2-1. The mapper is applied to each input key-value pair, producing an output key-value pair

As an example, consider a mapper whose purpose is to transform sentences into words. The input to this mapper would be strings that contain sentences, and the mapper's function would be to split the sentences into words and output the words (**Figure 2-2**).

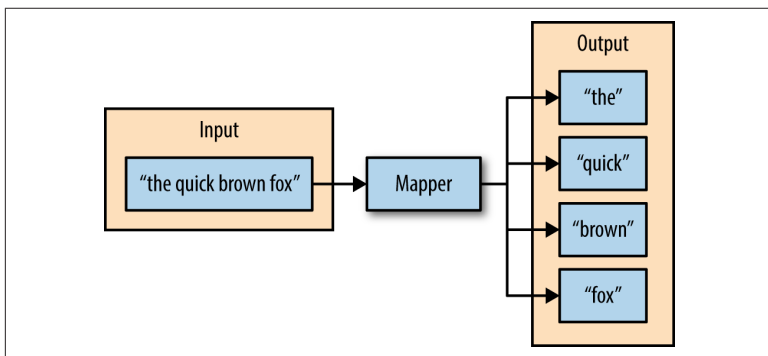


Figure 2-2. The input of the mapper is a string, and the function of the mapper is to split the input on spaces; the resulting output is the individual words from the mapper's input

Shuffle and Sort

The second phase of MapReduce is the shuffle and sort. As the mappers begin completing, the intermediate outputs from the map phase are moved to the reducers. This process of moving output from the mappers to the reducers is known as shuffling.

Shuffling is handled by a partition function, known as the partitioner. The partitioner is used to control the flow of key-value pairs from mappers to reducers. The partitioner is given the mapper's output key and the number of reducers, and returns the index of the intended reducer. The partitioner ensures that all of the values for the same key are sent to the same reducer. The default partitioner is hash-based. It computes a hash value of the mapper's output key and assigns a partition based on this result.

The final stage before the reducers start processing data is the sorting process. The intermediate keys and values for each partition are sorted by the Hadoop framework before being presented to the reducer.

Reduce

The third phase of MapReduce is the reduce phase. Within the reducer phase, an iterator of values is provided to a function known as the reducer. The iterator of values is a nonunique set of values for each unique key from the output of the map phase. The reducer aggregates the values for each unique key and produces zero or more output key-value pairs (Figure 2-3).




Figure 2-3. The reducer iterates over the input values, producing an output key-value pair

As an example, consider a reducer whose purpose is to sum all of the values for a key. The input to this reducer is an iterator of all of the values for a key, and the reducer sums all of the values. The reducer then outputs a key-value pair that contains the input key and the sum of the input key values (Figure 2-4).




Figure 2-4. This reducer sums the values for the keys “cat” and “mouse”

The next section describes a simple MapReduce application and its implementation in Python.

Hadoop Streaming

Hadoop streaming is a utility that comes packaged with the Hadoop distribution and allows MapReduce jobs to be created with any executable as the mapper and/or the reducer. The Hadoop streaming utility enables Python, shell scripts, or any other language to be used as a mapper, reducer, or both.

How It Works

The mapper and reducer are both executables that read input, line by line, from the standard input (stdin), and write output to the standard output (stdout). The Hadoop streaming utility creates a MapReduce job, submits the job to the cluster, and monitors its progress until it is complete.

When the mapper is initialized, each map task launches the specified executable as a separate process. The mapper reads the input file and presents each line to the executable via stdin. After the executable processes each line of input, the mapper collects the output from stdout and converts each line to a key-value pair. The key consists of the part of the line before the first tab character, and the value consists of the part of the line after the first tab character. If a line contains no tab character, the entire line is considered the key and the value is null.

When the reducer is initialized, each reduce task launches the specified executable as a separate process. The reducer converts the input key-value pair to lines that are presented to the executable via stdin. The reducer collects the executables result from stdout and converts each line to a key-value pair. Similar to the mapper, the executable specifies key-value pairs by separating the key and value by a tab character.

A Python Example

To demonstrate how the Hadoop streaming utility can run Python as a MapReduce application on a Hadoop cluster, the WordCount application can be implemented as two Python programs: *mapper.py* and *reducer.py*.

mapper.py is the Python program that implements the logic in the map phase of WordCount. It reads data from stdin, splits the lines into words, and outputs each word with its intermediate count to stdout. The code in [Example 2-1](#) implements the logic in *mapper.py*.

Example 2-1. python/MapReduce/HadoopStreaming/mapper.py

```
#!/usr/bin/env python

import sys

# Read each line from stdin
```

```
for line in sys.stdin:

    # Get the words in each line
    words = line.split()

    # Generate the count for each word
    for word in words:

        # Write the key-value pair to stdout to be processed by
        # the reducer.
        # The key is anything before the first tab character and the
        # value is anything after the first tab character.
        print '{0}\t{1}'.format(word, 1)
```

reducer.py is the Python program that implements the logic in the reduce phase of WordCount. It reads the results of *mapper.py* from stdin, sums the occurrences of each word, and writes the result to stdout. The code in [Example 2-2](#) implements the logic in *reducer.py*.

Example 2-2. python/MapReduce/HadoopStreaming/reducer.py

```
#!/usr/bin/env python

import sys

curr_word = None
curr_count = 0

# Process each key-value pair from the mapper
for line in sys.stdin:

    # Get the key and value from the current line
    word, count = line.split('\t')

    # Convert the count to an int
    count = int(count)

    # If the current word is the same as the previous word,
    # increment its count, otherwise print the words count
    # to stdout
    if word == curr_word:
        curr_count += count
    else:

        # Write word and its number of occurrences as a key-value
        # pair to stdout
        if curr_word:
            print '{0}\t{1}'.format(curr_word, curr_count)

        curr_word = word
```

```

    curr_count = count

# Output the count for the last word
if curr_word == word:
    print '{0}\t{1}'.format(curr_word, curr_count)

```

Before attempting to execute the code, ensure that the *mapper.py* and *reducer.py* files have execution permission. The following command will enable this for both files:

```
$ chmod a+x mapper.py reducer.py
```

Also ensure that the first line of each file contains the proper path to Python. This line enables *mapper.py* and *reducer.py* to execute as standalone executables. The value `#!/usr/bin/env python` should work for most systems, but if it does not, replace `/usr/bin/env python` with the path to the Python executable on your system.

To test the Python programs locally before running them as a Map-Reduce job, they can be run from within the shell using the `echo` and `sort` commands. It is highly recommended to test all programs locally before running them across a Hadoop cluster.

```

$ echo 'jack be nimble jack be quick' | ./mapper.py
| sort -t 1 | ./reducer.py
be      2
jack    2
nimble  1
quick   1

```

Once the mapper and reducer programs are executing successfully against tests, they can be run as a MapReduce application using the Hadoop streaming utility. The command to run the Python programs *mapper.py* and *reducer.py* on a Hadoop cluster is as follows:

```

$ $HADOOP_HOME/bin/hadoop jar
  $HADOOP_HOME/mapred/contrib/streaming/hadoop-streaming*.jar \
  -files mapper.py,reducer.py \
  -mapper mapper.py \
  -reducer reducer.py \
  -input /user/hduser/input.txt -output /user/hduser/output

```

The options used with the Hadoop streaming utility are listed in [Table 2-1](#).

Table 2-1. Options for Hadoop streaming

Option	Description
-files	A command-separated list of files to be copied to the MapReduce cluster
-mapper	The command to be run as the mapper
-reducer	The command to be run as the reducer
-input	The DFS input path for the Map step
-output	The DFS output directory for the Reduce step

mrjob

mrjob is a Python MapReduce library, created by Yelp, that wraps Hadoop streaming, allowing MapReduce applications to be written in a more Pythonic manner. mrjob enables multistep MapReduce jobs to be written in pure Python. MapReduce jobs written with mrjob can be tested locally, run on a Hadoop cluster, or run in the cloud using Amazon Elastic MapReduce (EMR).

Writing MapReduce applications with mrjob has many benefits:

- mrjob is currently a very actively developed framework with multiple commits every week.
- mrjob has extensive documentation, more than any other framework or library that supports Python on Hadoop.
- mrjob applications can be executed and tested without having Hadoop installed, enabling development and testing before deploying to a Hadoop cluster.
- mrjob allows MapReduce applications to be written in a single class, instead of writing separate programs for the mapper and reducer.

While mrjob is a great solution, it does have its drawbacks. mrjob is simplified, so it doesn't give the same level of access to Hadoop that other APIs offer. mrjob does not use typedbytes, so other libraries may be faster.

Installation

The installation of mrjob is simple; it can be installed with `pip` by using the following command:

```
$ pip install mrjob
```

Or it can be installed from source (a git clone):

```
$ python setup.py install
```

WordCount in mrjob

Example 2-3 uses mrjob to implement the WordCount algorithm.

Example 2-3. python/MapReduce/mrjob/word_count.py

```
from mrjob.job import MRJob

class MRWordCount(MRJob):

    def mapper(self, _, line):
        for word in line.split():
            yield(word, 1)

    def reducer(self, word, counts):
        yield(word, sum(counts))

if __name__ == '__main__':
    MRWordCount.run()
```

To run the mrjob locally, the only thing needed is a body of text. To run the job locally and count the frequency of words within a file named *input.txt*, use the following command:

```
$ python word_count.py input.txt
```

The output depends on the contents of the input file, but should look similar to **Example 2-4**.

Example 2-4. Output from word_count.py

```
"be"      2
"jack"    2
"nimble"  1
"quick"   1
```

What Is Happening

The MapReduce job is defined as the class, MRWordCount. Within the mrjob library, the class that inherits from MRJob contains the methods that define the steps of the MapReduce job. The steps within an mrjob application are mapper, combiner, and reducer. The class inheriting MRJob only needs to define one of these steps.

The `mapper()` method defines the mapper for the MapReduce job. It takes key and value as arguments and yields tuples of (output_key, output_value). In the WordCount example (Example 2-4), the mapper ignored the input key and split the input value to produce words and counts.

The `combiner()` method defines the combiner for the MapReduce job. The combiner is a process that runs after the mapper and before the reducer. It receives, as input, all of the data emitted by the mapper, and the output of the combiner is sent to the reducer. The combiner's input is a key, which was yielded by the mapper, and a value, which is a generator that yields all values yielded by one mapper that corresponds to the key. The combiner yields tuples of (output_key, output_value) as output.

The `reducer()` method defines the reducer for the MapReduce job. It takes a key and an iterator of values as arguments and yields tuples of (output_key, output_value). In Example 2-4, the reducer sums the value for each key, which represents the frequency of words in the input.

The final component of a MapReduce job written with the `mrjob` library is the two lines at the end of the file:

```
if __name__ == '__main__':  
    MRWordCount.run()
```

These lines enable the execution of `mrjob`; without them, the application will not work.

Executing mrjob

Executing a MapReduce application with `mrjob` is similar to executing any other Python program. The command line must contain the name of the `mrjob` application and the input file:

```
$ python mr_job.py input.txt
```

By default, `mrjob` writes output to stdout.

Multiple files can be passed to `mrjob` as inputs by specifying the filenames on the command line:

```
$ python mr_job.py input1.txt input2.txt input3.txt
```

`mrjob` can also handle input via stdin:

```
$ python mr_job.py < input.txt
```


By default, mrjob runs locally, allowing code to be developed and debugged before being submitted to a Hadoop cluster.

To change how the job is run, specify the `-r/--runner` option. **Table 2-2** contains a description of the valid choices for the runner options.

Table 2-2. mrjob runner choices

<code>-r inline</code>	(Default) Run in a single Python process
<code>-r local</code>	Run locally in a few subprocesses simulating some Hadoop features
<code>-r hadoop</code>	Run on a Hadoop cluster
<code>-r emr</code>	Run on Amazon Elastic Map Reduce (EMR)

Using the runner option allows the mrjob program to be run on a Hadoop cluster, with input being specified from HDFS:

```
$ python mr_job.py -r hadoop hdfs://input/input.txt
```

mrjob also allows applications to be run on EMR directly from the command line:

```
$ python mr_job.py -r emr s3://input-bucket/input.txt
```

Top Salaries

Example 2-5 uses mrjob to compute employee top annual salaries and gross pay. The dataset used is the **salary information** from the city of Baltimore for 2014.

Example 2-5. python/MapReduce/mrjob/top_salary.py

```
from mrjob.job import MRJob
from mrjob.step import MRStep
import csv

cols = 'Name,JobTitle,AgencyID,Agency,HireDate,AnnualSalary,Gross
Pay'.split(',')

class salarymax(MRJob):

    def mapper(self, _, line):
        # Convert each line into a dictionary
        row = dict(zip(cols, [ a.strip() for a in
            csv.reader([line]).next()])))

        # Yield the salary
```

```

yield 'salary', (float(row['AnnualSalary'])[1:]), line)

# Yield the gross pay
try:
    yield 'gross', (float(row['GrossPay'])[1:]), line)
except ValueError:
    self.increment_counter('warn', 'missing gross', 1)

def reducer(self, key, values):
    topten = []

    # For 'salary' and 'gross' compute the top 10
    for p in values:
        topten.append(p)
        topten.sort()
        topten = topten[-10:]

    for p in topten:
        yield key, p

combiner = reducer

if __name__ == '__main__':
    salarymax.run()

```

Use the following command to execute the MapReduce job on Hadoop:

```
$ python top_salary.py -r hadoop hdfs:///user/hduser/input/
salaries.csv
```

Chapter Summary

This chapter introduced the MapReduce programming model and described how data flows through the different phases of the model. Hadoop Streaming and mrjob were then used to highlight how MapReduce jobs can be written in Python.

Pig and Python

Pig is composed of two major parts: a high-level data flow language called Pig Latin, and an engine that parses, optimizes, and executes the Pig Latin scripts as a series of MapReduce jobs that are run on a Hadoop cluster. Compared to Java MapReduce, Pig is easier to write, understand, and maintain because it is a data transformation language that allows the processing of data to be described as a sequence of transformations. Pig is also highly extensible through the use of the User Defined Functions (UDFs) which allow custom processing to be written in many languages, such as Python.

An example of a Pig application is the Extract, Transform, Load (ETL) process that describes how an application extracts data from a data source, transforms the data for querying and analysis purposes, and loads the result onto a target data store. Once Pig loads the data, it can perform projections, iterations, and other transformations. UDFs enable more complex algorithms to be applied during the transformation phase. After the data is done being processed by Pig, it can be stored back in HDFS.

This chapter begins with an example Pig script. Pig and Pig Latin are then introduced and described in detail with examples. The chapter concludes with an explanation of how Pig's core features can be extended through the use of Python.

WordCount in Pig

Example 3-1 implements the WordCount algorithm in Pig. It assumes that a data file, *input.txt*, is loaded in HDFS under */user/hduser/input*, and output will be placed in HDFS under */user/hduser/output*.

Example 3-1. pig/wordcount.pig

```
%default INPUT '/user/hduser/input/input.txt';
%default OUTPUT '/user/hduser/output';

-- Load the data from the file system into the relation records
records = LOAD '$INPUT';

-- Split each line of text and eliminate nesting
terms = FOREACH records GENERATE FLATTEN(TOKENIZE((chararray) $0))
AS word;

-- Group similar terms
grouped_terms = GROUP terms BY word;

-- Count the number of tuples in each group
word_counts = FOREACH grouped_terms GENERATE COUNT(terms), group;

-- Store the result
STORE word_counts INTO '$OUTPUT';
```

To execute the Pig script, simply call Pig from the command line and pass it the name of the script to run:

```
$ pig wordcount.pig
```

While the job is running, a lot of text will be printed to the console. Once the job is complete, a success message, similar to the one below, will be displayed:

```
2015-09-26 14:15:10,030 [main] INFO org.apache.pig.back-
end.hadoop.executionengine.mapReduceLayer.MapReduceLauncher -
Success!
2015-09-26 14:15:10,049 [main] INFO org.apache.pig.Main - Pig
script completed in 18 seconds and 514 milliseconds (18514 ms)
```

The results of the *wordcount.pig* script are displayed in **Example 3-2** and can be found in HDFS under */user/hduser/output/pig_wordcount/part-r-000000*.

Example 3-2. */user/hduser/output/pig_wordcount/part-r-00000*

```
2 be
1 the
3 jack
1 over
1 quick
1 jumped
1 nimble
1 candlestick
```

WordCount in Detail

This section describes each Pig Latin statement in the *wordcount.pig* script.

The first statement loads data from the filesystem and stores it in the relation records:

```
records = LOAD '/user/hduser/input/input.txt';
```

The second statement splits each line of text using the TOKENIZE function and eliminates nesting using the FLATTEN operator:

```
terms = FOREACH records GENERATE FLATTEN(TOKENIZE((chararray)
$0)) AS word;
```

The third statement uses the GROUP operator to group the tuples that have the same field:

```
grouped_terms = GROUP terms BY word;
```

The fourth statement iterates over all of the terms in each bag and uses the COUNT function to return the sum:

```
word_counts = FOREACH grouped_terms GENERATE COUNT(terms),
group;
```

The fifth and final statement stores the results in HDFS:

```
STORE word_counts INTO '/user/hduser/output/pig_wordcount'
```

Running Pig

Pig contains multiple modes that can be specified to configure how Pig scripts and Pig statements will be executed.

Execution Modes

Pig has two execution modes: local and MapReduce.

Running Pig in local mode only requires a single machine. Pig will run on the local host and access the local filesystem. To run Pig in local mode, use the `-x local` flag:

```
$ pig -x local ...
```

Running Pig in MapReduce mode requires access to a Hadoop cluster. MapReduce mode executes Pig statements and jobs on the cluster and accesses HDFS. To run Pig in MapReduce mode, simply call Pig from the command line or use the `-x mapreduce` flag:

```
$ pig ...  
or  
$ pig -x mapreduce ...
```

Interactive Mode

Pig can be run interactively in the Grunt shell. To invoke the Grunt shell, simply call Pig from the command line and specify the desired execution mode. The following example starts the Grunt shell in local mode:

```
pig -x local  
...  
grunt>
```

Once the Grunt shell is initialized, Pig Latin statements can be entered and executed in an interactive manner. Running Pig interactively is a great way to learn Pig.

The following example reads `/etc/passwd` and displays the usernames from within the Grunt shell:

```
grunt> A = LOAD '/etc/passwd' using PigStorage(':');  
grunt> B = FOREACH A GENERATE $0 as username;  
grunt> DUMP B;
```

Batch Mode

Batch mode allows Pig to execute Pig scripts in local or MapReduce mode.

The Pig Latin statements in [Example 3-3](#) read a file named `passwd` and use the STORE operator to store the results in a directory called `user_id.out`. Before executing this script, ensure that `/etc/passwd` is copied to the current working directory if Pig will be run in local mode, or to HDFS if Pig will be executed in MapReduce mode.

Example 3-3. *pig/user_id.pig*

```
A = LOAD 'passwd' USING PigStorage(':');
B = FOREACH A GENERATE $0 as username;
STORE B INTO 'user_id.out';
```

Use the following command to execute the *user_id.pig* script on the local machine:

```
$ pig -x local user_id.pig
```

Pig Latin

This section describes the basic concepts of the Pig Latin language, allowing those new to the language to understand and write basic Pig scripts. For a more comprehensive overview of the language, visit the [Pig online documentation](#).

All of the examples in this section load and process data from the tab-delimited file, *resources/students* ([Example 3-4](#)).

Example 3-4. *resources/students*

john	21	3.89
sally	19	2.56
alice	22	3.76
doug	19	1.98
susan	26	3.25

Statements

Statements are the basic constructs used to process data in Pig. Each statement is an operator that takes a relation as an input, performs a transformation on that relation, and produces a relation as an output. Statements can span multiple lines, but all statements must end with a semicolon (;).

The general form of each Pig script is as follows:

1. A LOAD statement that reads the data from the filesystem
2. One or more statements to transform the data
3. A DUMP or STORE statement to view or store the results, respectively

Loading Data

The LOAD operator is used to load data from the system into Pig. The format of the LOAD operator is as follows:

```
LOAD 'data' [USING function] [AS schema];
```

Where 'data' is the name of the file or directory, in quotes, to be loaded. If a directory name is not specified, all of the files within the directory are loaded.

The USING keyword is optional and is used to specify a function to parse the incoming data. If the USING keyword is omitted, the default loading function, PigStorage, is used. The default delimiter is the tab character ('\t').

The AS keyword allows a schema to be defined for the data being loaded. Schemas enable names and datatypes to be declared for individual fields. The following example defines a schema for the data being loaded from the file *input.txt*. If no schema is defined, the fields are not named and default to type bytearray.

```
A = LOAD 'students' AS (name:chararray, age:int);

DUMP A;
(john,21,3.89)
(sally,19,2.56)
(alice,22,3.76)
(doug,19,1.98)
(susan,26,3.25)
```

Transforming Data

Pig contains many operators that enable complex transforming of data. The most common operators are FILTER, FOREACH, and GROUP.

FILTER

The FILTER operator works on tuples or rows of data. It selects tuples from a relation based on a condition.

The following examples use the relation A that contains student data:

```
A = LOAD 'students' AS (name:chararray, age:int, gpa:float);

DUMP A;
(john,21,3.89)
(sally,19,2.56)
(alice,22,3.76)
```



```
(doug,19,1.98)
(susan,26,3.25)
```

The following example filters out any students under the age of 20, and stores the results in a relation R:

```
R = FILTER A BY age >= 20;
```

```
DUMP R;
(john,21,3.89)
(alice,22,3.76)
(susan,26,3.25)
```

Condition statements can use the AND, OR, and NOT operators to create more complex FILTER statements. The following example filters out any students with an age less than 20 or a GPA less than or equal to 3.5, and stores the results in a relation R:

```
R = FILTER A BY (age >= 20) AND (gpa > 3.5)
```

```
DUMP R;
(john,21,3.89)
(alice,22,3.76)
```

FOREACH

While the FILTER operator works on rows of data, the FOREACH operator works on columns of data and is similar to the SELECT statement in SQL.

The following example uses the asterisk (*) to project all of the fields from relation A onto relation X:

```
R = FOREACH A GENERATE *;
```

```
DUMP R;
(john,21,3.89)
(sally,19,2.56)
(alice,22,3.76)
(doug,19,1.98)
(susan,26,3.25)
```

The following example uses field names to project the age and gpa columns from relation A onto relation X:

```
R = FOREACH A GENERATE age, gpa;
```

```
DUMP R;
(21,3.89)
(19,2.56)
(22,3.76)
```

```
(19,1.98)
(26,3.25)
```

GROUP

The GROUP operator groups together tuples that have the same group key into one or more relations.

The following example groups the student data by age and stores the result into relation B:

```
B = GROUP A BY age;

DUMP B;
(19, {(doug,19,1.98),(sally,19,2.56)})
(21, {(john,21,3.89)})
(22, {(alice,22,3.76)})
(26, {(susan,26,3.25)})
```

The result of a GROUP operation is a relation that has one tuple per group. This tuple has two fields: the first field is named group and is of the type of the grouped key; the second field is a bag that takes the name of the original relation. To clarify the structure of relation B, the DESCRIBE and ILLUSTRATE operations can be used:

```
DESCRIBE B;
B: {group: int, A: {(name: chararray, age: int, gpa: float)}}
```

```
ILLUSTRATE B;
```

```
-----
| B | group:int | A:bag{:tuple(name:chararray,          |
|   |         | age:int,gpa:float)}                  |
-----
|   | 19       | {(sally, 19, 2.56), (doug, 19, 1.98)} |
-----
```

Using the FOREACH operator, the fields in the previous relation, B, can be referred to by names group and A:

```
C = FOREACH B GENERATE group, A.name;

DUMP C;
(19, {(doug),(sally)})
(21, {(john)})
(22, {(alice)})
(26, {(susan)})
```

Storing Data

The STORE operator is used to execute previous Pig statements and store the results on the filesystem. The format of the STORE operator is as follows:

```
STORE alias INTO 'directory' [USING function];
```

Where `alias` is the name of the relation to store, and `'directory'` is the name of the storage directory, in quotes. If the directory already exists, the STORE operation will fail. The output files will be named *part-nnnnn* and are written to the specified directory.

The USING keyword is optional and is used to specify a function to store the data. If the USING keyword is omitted, the default storage function, `PigStorage`, is used. The following example specifies the `PigStorage` function to store a file with pipe-delimited fields:

```
A = LOAD 'students' AS (name:chararray, age:int, gpa:float);
```

```
DUMP A;
(john,21,3.89)
(sally,19,2.56)
(alice,22,3.76)
(doug,19,1.98)
(susan,26,3.25)
```

```
STORE A INTO 'output' USING PigStorage('|');
```

```
CAT output;
john|21|3.89
sally|19|2.56
alice|22|3.76
doug|19|1.98
susan|26|3.25
```

The provided Pig Latin statements are great general-purpose computing constructs, but are not capable of expressing complex algorithms. The next section describes how to extend the functionality of Pig with Python.

Extending Pig with Python

Pig provides extensive support for custom processing through User Defined Functions (UDFs). Pig currently supports UDFs in six languages: Java, Jython, Python, JavaScript, Ruby, and Groovy.

When Pig executes, it automatically detects the usage of a UDF. To run Python UDFs, Pig invokes the Python command line and streams data in and out of it.

Registering a UDF

Before a Python UDF can be used in a Pig script, it must be registered so Pig knows where to look when the UDF is called. To register a Python UDF file, use Pig's REGISTER statement:

```
REGISTER 'udfs/myudf.py' USING streaming_python AS my_udf;
```

Once the UDF is registered, it can be called from within the Pig script:

```
relation = FOREACH data GENERATE my_udf.function(field);
```

In this example the UDF, referenced as `my_udf`, contains a function called `function`.

A Simple Python UDF

A simple Python UDF, located in `pig/udfs/my_first_udf.py`, that returns the integer value 1 each time it is called, is shown in [Example 3-5](#).

Example 3-5. pig/udfs/my_first_udf.py

```
from pig_util import outputSchema

@outputSchema('value:int')
def return_one():
    """
    Return the integer value 1
    """
    return 1
```

Some important things to note in this Python script are the `from` statement on the first line, and the output decorator, `@outputSchema` decorator, on the third line. These lines enable the Python UDF to define an alias and datatype for the data being returned from the UDF.

The Pig script in [Example 3-6](#) registers the Python UDF and calls the `return_one()` function in a `FOREACH` statement.

Example 3-6. pig/simple_udf.pig

```
REGISTER 'udfs/my_first_udf.py' USING streaming_python AS pyudfs;

A = LOAD './resources/input.txt';
B = FOREACH A GENERATE pyudfs.return_one();
DUMP B;
```

When the Pig script is executed, it will generate an integer value 1 for each line in the input file. Use the following command to execute the script (sample output is shown as well):

```
$ pig -x local simple_udf.pig
...
(1)
(1)
(1)
```

String Manipulation

Python UDFs are an easy way of extending Pig's functionality and an easy way to transform and process data.

The Python UDF in [Example 3-7](#) contains two functions: `reverse()` and `num_chars()`. The `reverse()` function takes in a chararray and returns the chararray in reverse order. The `num_chars()` function takes in a chararray and returns the number of characters in the chararray.

Example 3-7. pig/udfs/string_funcs.py

```
from pig_util import outputSchema

@outputSchema('word:chararray')
def reverse(word):
    """
    Return the reverse text of the provided word
    """
    return word[::-1]

@outputSchema('length:int')
def num_chars(word):
    """
    Return the length of the provided word
    """
    return len(word)
```

The Pig script in [Example 3-8](#) loads a text file and applies the `reverse()` and `num_chars()` Python functions to each unique word.

Example 3-8. pig/playing_with_words.pig

```
REGISTER 'udfs/string_funcs.py' USING streaming_python AS
string_udf;

-- Load the data from the file system
records = LOAD './resources/input.txt';

-- Split each line of text and eliminate nesting
terms = FOREACH records GENERATE FLATTEN(TOKENIZE((chararray) $0))
AS word;

-- Group similar terms
grouped_terms = GROUP terms BY word;

-- Count the number of tuples in each group
unique_terms = FOREACH grouped_terms GENERATE group as word;

-- Calculate the number of characters in each term
term_length = FOREACH unique_terms GENERATE word,
string_udf.num_chars(word) as length;

-- Display the terms and their length
DUMP term_length;

-- Reverse each word
reverse_terms = FOREACH unique_terms GENERATE word,
string_udf.reverse(word) as reverse_word;

-- Display the terms and the reverse terms
DUMP reverse_terms;
```

Use the following command to execute the script (sample output shown):

```
$ pig -x local playing_with_words.pig
...
(be,2)
(the,3)
(jack,4)
(over,4)
(quick,5)
(jumped,6)
(nimble,6)
(candlestick,11)
...
(be,eb)
```

```
(the,eht)
(jack,kcaj)
(over,revo)
(quick,kciuq)
(jumped,depnuj)
(nimble,elbmim)
(candlestick,kcitseldnac)
```

Most Recent Movies

The following example uses movie data from the [groupLens datasets](#) and external libraries to calculate the 10 most recent movies.

The Python UDF in [Example 3-9](#) contains two functions: `parse_title()` and `days_since_release()`. The `parse_title()` function uses Python's regular expression module to remove the release year from a movie's title. The `days_since_release()` function uses the `datetime` module to calculate the number of days between the current day and a movie's release date.

Example 3-9. `pig/udfs/movies_udf.py`

```
from pig_util import outputSchema
from datetime import datetime
import re

@outputSchema('title:chararray')
def parse_title(title):
    """
    Return the title without the year
    """
    return re.sub(r'\s*\(\d{4}\)', '', title)

@outputSchema('days_since_release:int')
def days_since_release(date):
    """
    Calculate the number of days since the titles release
    """
    if date is None:
        return None

    today = datetime.today()
    release_date = datetime.strptime(date, '%d-%b-%Y')
    delta = today - release_date
    return delta.days
```

The Pig script in **Example 3-10** uses the Python UDFs to determine the 10 most recent movies.

Example 3-10. pig/recent_movies.pig

```
REGISTER 'udfs/movies_udf.py' USING streaming_python AS movies_udf;

-- Load the data from the file system
records = LOAD '../resources/movies' USING PigStorage('|')
          AS (id:int, title:chararray, release_date:chararray);

-- Parse the titles and determine how many days since the release
date
titles = FOREACH records GENERATE movies_udf.parse_title(title),
movies_udf.days_since_release(release_date);

-- Order the movies by the time since release
most_recent = ORDER titles BY days_since_release ASC;

-- Get the ten most recent movies
top_ten = LIMIT most_recent 10;

-- Display the top ten most recent movies
DUMP top_ten;
```

The following command is used to execute the script (sample output shown):

```
$ pig -x local recent_movies.pig
...
(unknown,)
(Apt Pupil,6183)
(Mighty, The,6197)
(City of Angels,6386)
(Big One, The,6393)
(Lost in Space,6393)
(Mercury Rising,6393)
(Spanish Prisoner, The,6393)
(Hana-bi,6400)
(Object of My Affection, The,6400)
```

Chapter Summary

This chapter introduced and Pig and Pig Latin. It described the basic concepts of Pig Latin, allowing simple Pig scripts to be created and executed. It also introduced how to extend the functionality of Pig Latin with Python UDFs.

Spark with Python

Spark is a cluster computing framework that uses in-memory primitives to enable programs to run up to a hundred times faster than Hadoop MapReduce applications. Spark applications consist of a driver program that controls the execution of parallel operations across a cluster. The main programming abstraction provided by Spark is known as Resilient Distributed Datasets (RDDs). RDDs are collections of elements partitioned across the nodes of the cluster that can be operated on in parallel.

Spark was created to run on many platforms and be developed in many languages. Currently, Spark can run on Hadoop 1.0, Hadoop 2.0, Apache Mesos, or a standalone Spark cluster. Spark also natively supports Scala, Java, Python, and R. In addition to these features, Spark can be used interactively from a command-line shell.

This chapter begins with an example Spark script. PySpark is then introduced, and RDDs are described in detail with examples. The chapter concludes with example Spark programs written in Python.

WordCount in PySpark

The code in **Example 4-1** implements the WordCount algorithm in PySpark. It assumes that a data file, *input.txt*, is loaded in HDFS under */user/hduser/input*, and the output will be placed in HDFS under */user/hduser/output*.

Example 4-1. *python/Spark/word_count.py*

```
from pyspark import SparkContext

def main():

    sc = SparkContext(appName='SparkWordCount')

    input_file = sc.textFile('/user/hduser/input/input.txt')
    counts = input_file.flatMap(lambda line: line.split()) \
        .map(lambda word: (word, 1)) \
        .reduceByKey(lambda a, b: a + b)
    counts.saveAsTextFile('/user/hduser/output')

    sc.stop()

if __name__ == '__main__':
    main()
```

To execute the Spark application, pass the name of the file to the spark-submit script:

```
$ spark-submit --master local word_count.py
```

While the job is running, a lot of text will be printed to the console. The results of a *word_count.py* Spark script are displayed in [Example 4-2](#) and can be found in HDFS under */user/hduser/output/part-00000*.

Example 4-2. */user/hduser/output/part-00000*

```
(u'be', 2)
(u'jumped', 1)
(u'over', 1)
(u'candlestick', 1)
(u'nimble', 1)
(u'jack', 3)
(u'quick', 1)
(u'the', 1)
```

WordCount Described

This section describes the transformations being applied in the *word_count.py* Spark script.

The first statement creates a *SparkContext* object. This object tells Spark how and where to access a cluster:

```
sc = SparkContext(appName='SparkWordCount')
```

The second statement uses the `SparkContext` to load a file from HDFS and store it in the variable `input_file`:

```
input_file = sc.textFile('/user/hduser/input/input.txt')
```

The third statement performs multiple transformations on the input data. Spark automatically parallelizes these transformations to run across multiple machines:

```
counts = input_file.flatMap(lambda line: line.split()) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
```

The fourth statement stores the results to HDFS:

```
counts.saveAsTextFile('/user/hduser/output')
```

The fifth statement shuts down the `SparkContext`:

```
sc.stop()
```

PySpark

PySpark is Spark's Python API. PySpark allows Spark applications to be created from an interactive shell or from Python programs.

Before executing any code within Spark, the application must create a `SparkContext` object. The `SparkContext` object tells Spark how and where to access a cluster. The `master` property is a cluster URL that determines where the Spark application will run. The most common values for `master` are:

`local`

Run Spark with one worker thread.

`local[n]`

Run Spark with n worker threads.

`spark://HOST:PORT`

Connect to a Spark standalone cluster.

`mesos://HOST:PORT`

Connect to a Mesos cluster.

Interactive Shell

In the Spark shell, the `SparkContext` is created when the shell launches. The `SparkContext` is held in the variable `sc`. The master

for the interactive shell can be set by using the `--master` argument when the shell is launched. To start an interactive shell, run the `pyspark` command:

```
$ pyspark --master local[4]
...
Welcome to

      /--_/ _ _ _ _ _ \_/_/
     /  \  \  \  \  \  \  \
    / _ /  .  \ \ _ _ \ / \ \
   / _ /                                     version 1.5.0
  / _ /

Using Python version 2.7.10 (default, Jul 13 2015 12:05:58)
SparkContext available as sc, HiveContext available as sqlCon-
text.
>>>
```

For a complete list of options, run `pyspark --help`.

Self-Contained Applications

Self-contained applications must first create a `SparkContext` object before using any Spark methods. The master can be set when the `SparkContext()` method is called:

```
sc = SparkContext(master='local[4]')
```

To execute self-contained applications, they must be submitted to the `spark-submit` script. The `spark-submit` script contains many options; to see a complete listing, run `spark-submit --help` from the command line:

```
$ spark-submit --master local spark_app.py
```

Resilient Distributed Datasets (RDDs)

Resilient Distributed Datasets (RDDs) are the fundamental programming abstraction in Spark. RDDs are immutable collections of data, partitioned across machines, that enable operations to be performed on elements in parallel. RDDs can be constructed in multiple ways: by parallelizing existing Python collections, by referencing files in an external storage system such as HDFS, or by applying transformations to existing RDDs.

Creating RDDs from Collections

RDDs can be created from a Python collection by calling the `SparkContext.parallelize()` method. The elements of the collection are copied to form a distributed dataset that can be operated on in parallel. The following example creates a parallelized collection from a Python list:

```
>>> data = [1, 2, 3, 4, 5]
>>> rdd = sc.parallelize(data)
>>> rdd.glom().collect()
...
[[1, 2, 3, 4, 5]]
```

The `RDD.glom()` method returns a list of all of the elements within each partition, and the `RDD.collect()` method brings all the elements to the driver node. The result, `[[1, 2, 3, 4, 5]]`, is the original collection within a list.

To specify the number of partitions an RDD should be created with, a second argument can be passed to the `parallelize()` method. The following example creates an RDD from the same Python collection in the previous example, except this time four partitions are created:

```
>>> rdd = sc.parallelize(data, 4)
>>> rdd.glom().collect()
...
[[1], [2], [3], [4, 5]]
```

Using the `glom()` and `collect()` methods, the RDD created in this example contains four inner lists: `[1]`, `[2]`, `[3]`, and `[4, 5]`. The number of inner lists represents the number of partitions within the RDD.

Creating RDDs from External Sources

RDDs can also be created from files using the `SparkContext.textFile()` method. Spark can read files residing on the local filesystem, any storage source supported by Hadoop, Amazon S3, and so on. Spark supports text files, SequenceFiles, any other Hadoop Input-Format, directories, compressed files, and wildcards, e.g., `my/directory/*.txt`. The following example creates a distributed dataset from a file located on the local filesystem:

```
>>> distFile = sc.textFile('data.txt')
>>> distFile.glom().collect()
...
[[u'jack be nimble', u'jack be quick', u'jack jumped over the
candlestick']]
```

As before, the `glom()` and `collect()` methods allow the RDD to be displayed in its partitions. This result shows that `distFile` only has a single partition.

Similar to the `parallelize()` method, the `textFile()` method takes a second parameter that specifies the number of partitions to create. The following example creates an RDD with three partitions from the input file:

```
>>> distFile = sc.textFile('data.txt', 3)
>>> distFile.glom().collect()
...
[[u'jack be nimble', u'jack be quick'], [u'jack jumped over
the candlestick'], []]
```

RDD Operations

RDDs support two types of operations: transformations and actions. Transformations create new datasets from existing ones, and actions run a computation on the dataset and return results to the driver program.

Transformations are lazy: that is, their results are not computed immediately. Instead, Spark remembers all of the transformations applied to a base dataset. Transformations are computed when an action requires a result to be returned to the driver program. This allows Spark to operate efficiently and only transfer the results of the transformations before an action.

By default, transformations may be recomputed each time an action is performed on it. This allows Spark to efficiently utilize memory, but it may utilize more processing resources if the same transformations are constantly being processed. To ensure a transformation is only computed once, the resulting RDD can be persisted in memory using the `RDD.cache()` method.

RDD Workflow

The general workflow for working with RDDs is as follows:

1. Create an RDD from a data source.
2. Apply transformations to an RDD.
3. Apply actions to an RDD.

The following example uses this workflow to calculate the number of characters in a file:

```
>>> lines = sc.textFile('data.txt')
>>> line_lengths = lines.map(lambda x: len(x))
>>> document_length = line_lengths.reduce(lambda x,y: x+y)
>>> print document_length
59
```

The first statement creates an RDD from the external file *data.txt*. This file is not loaded at this point; the variable `lines` is just a pointer to the external source. The second statement performs a transformation on the base RDD by using the `map()` function to calculate the number of characters in each line. The variable `line_lengths` is not immediately computed due to the laziness of transformations. Finally, the `reduce()` method is called, which is an action. At this point, Spark divides the computations into tasks to run on separate machines. Each machine runs both the map and reduction on its local data, returning only the results to the driver program.

If the application were to use `line_lengths` again, it would be best to persist the result of the map transformation to ensure that the map would not be recomputed. The following line will save `line_lengths` into memory after the first time it is computed:

```
>>> line_lengths.persist()
```

Python Lambda Functions

Many of Spark's transformations and actions require function objects to be passed from the driver program to run on the cluster. The easiest way to define and pass a function is through the use of Python lambda functions.

Lambda functions are anonymous functions (i.e., they do not have a name) that are created at runtime. They can be used wherever function objects are required and are syntactically restricted to a single

expression. The following example shows a lambda function that returns the sum of its two arguments:

```
lambda a, b: a + b
```

Lambdas are defined by the keyword `lambda`, followed by a comma-separated list of arguments. A colon separates the function declaration from the function expression. The function expression is a single expression that produces a result for the provided arguments.

In the previous Spark example, the `map()` function uses the following lambda function:

```
lambda x: len(x)
```

This lambda has one argument and returns the length of the argument.

Transformations

Transformations create new datasets from existing ones. Lazy evaluation of transformation allows Spark to remember the set of transformations applied to the base RDD. This enables Spark to optimize the required calculations.

This section describes some of Spark's most common transformations. For a full listing of transformations, refer to Spark's [Python RDD API doc](#).

map. The `map(func)` function returns a new RDD by applying a function, *func*, to each element of the source. The following example multiplies each element of the source RDD by two:

```
>>> data = [1, 2, 3, 4, 5, 6]
>>> rdd = sc.parallelize(data)
>>> map_result = rdd.map(lambda x: x * 2)
>>> map_result.collect()
[2, 4, 6, 8, 10, 12]
```

filter. The `filter(func)` function returns a new RDD containing only the elements of the source that the supplied function returns as true. The following example returns only the even numbers from the source RDD:

```
>>> data = [1, 2, 3, 4, 5, 6]
>>> filter_result = rdd.filter(lambda x: x % 2 == 0)
>>> filter_result.collect()
[2, 4, 6]
```


distinct. The `distinct()` method returns a new RDD containing only the distinct elements from the source RDD. The following example returns the unique elements in a list:

```
>>> data = [1, 2, 3, 2, 4, 1]
>>> rdd = sc.parallelize(data)
>>> distinct_result = rdd.distinct()
>>> distinct_result.collect()
[4, 1, 2, 3]
```

flatMap. The `flatMap(func)` function is similar to the `map()` function, except it returns a flattened version of the results. For comparison, the following examples return the original element from the source RDD and its square. The example using the `map()` function returns the pairs as a list within a list:

```
>>> data = [1, 2, 3, 4]
>>> rdd = sc.parallelize(data)
>>> map = rdd.map(lambda x: [x, pow(x,2)])
>>> map.collect()
[[1, 1], [2, 4], [3, 9], [4, 16]]
```

While the `flatMap()` function concatenates the results, returning a single list:

```
>>> rdd = sc.parallelize()
>>> flat_map = rdd.flatMap(lambda x: [x, pow(x,2)])
>>> flat_map.collect()
[1, 1, 2, 4, 3, 9, 4, 16]
```

Actions

Actions cause Spark to compute transformations. After transforms are computed on the cluster, the result is returned to the driver program.

The following section describes some of Spark's most common actions. For a full listing of actions, refer to Spark's [Python RDD API doc](#).

reduce. The `reduce()` method aggregates elements in an RDD using a function, which takes two arguments and returns one. The function used in the reduce method is commutative and associative, ensuring that it can be correctly computed in parallel. The following example returns the product of all of the elements in the RDD:

```
>>> data = [1, 2, 3]
>>> rdd = sc.parallelize(data)
```

```
>>> rdd.reduce(lambda a, b: a * b)
6
```

take. The `take(n)` method returns an array with the first *n* elements of the RDD. The following example returns the first two elements of an RDD:

```
>>> data = [1, 2, 3]
>>> rdd = sc.parallelize(data)
>>> rdd.take(2)
[1, 2]
```

collect. The `collect()` method returns all of the elements of the RDD as an array. The following example returns all of the elements from an RDD:

```
>>> data = [1, 2, 3, 4, 5]
>>> rdd = sc.parallelize(data)
>>> rdd.collect()
[1, 2, 3, 4, 5]
```

It is important to note that calling `collect()` on large datasets could cause the driver to run out of memory. To inspect large RDDs, the `take()` and `collect()` methods can be used to inspect the top *n* elements of a large RDD. The following example will return the first 100 elements of the RDD to the driver:

```
>>> rdd.take(100).collect()
```

takeOrdered. The `takeOrdered(n, key=func)` method returns the first *n* elements of the RDD, in their natural order, or as specified by the function *func*. The following example returns the first four elements of the RDD in descending order:

```
>>> data = [6,1,5,2,4,3]
>>> rdd = sc.parallelize(data)
>>> rdd.takeOrdered(4, lambda s: -s)
[6, 5, 4, 3]
```

Text Search with PySpark

The text search program searches for movie titles that match a given string ([Example 4-3](#)). The movie data is from the [groupLens](#) datasets; the application expects this to be stored in HDFS under `/user/hduser/input/movies`.

Example 4-3. python/Spark/text_search.py

```
from pyspark import SparkContext
import re
import sys

def main():

    # Insure a search term was supplied at the command line
    if len(sys.argv) != 2:
        sys.stderr.write('Usage: {} <search_term>'.format(sys.argv[0]))
        sys.exit()

    # Create the SparkContext
    sc = SparkContext(appName='SparkWordCount')

    # Broadcast the requested term
    requested_movie = sc.broadcast(sys.argv[1])

    # Load the input file
    source_file = sc.textFile('/user/hduser/input/movies')

    # Get the movie title from the second fields
    titles = source_file.map(lambda line: line.split('|')[1])

    # Create a map of the normalized title to the raw title
    normalized_title = titles.map(lambda title: (re.sub(r'\s*\s*\s*\s*', ' ', title).lower(), title))

    # Find all movies matching the requested_movie
    matches = normalized_title.filter(lambda x: requested_movie.value in x[0])

    # Collect all the matching titles
    matching_titles = matches.map(lambda x: x[1]).distinct().collect()

    # Display the result
    print '{} Matching titles found: {}'.format(len(matching_titles))
    for title in matching_titles:
        print title

    sc.stop()

if __name__ == '__main__':
    main()
```

The Spark application can be executed by passing to the spark-submit script the name of the program, *text_search.py*, and the term

for which to search. A sample run of the application can be seen here:

```
$ spark-submit text_search.py gold
...
6 Matching titles found:
GoldenEye (1995)
On Golden Pond (1981)
Ulee's Gold (1997)
City Slickers II: The Legend of Curly's Gold (1994)
Golden Earrings (1947)
Gold Diggers: The Secret of Bear Mountain (1995)
...
```

Since computing the transformations can be a costly operation, Spark can cache the results of the `normalized_titles` to memory to speed up future searches. From the example above, to load the `normalized_titles` into memory, use the `cache()` method:

```
normalized_title.cache()
```

Chapter Summary

This chapter introduced Spark and PySpark. It described Spark's main programming abstraction, RDDs, with many examples of dataset transformations. This chapter also contained a Spark application that returned movie titles that matched a given string.

Workflow Management with Python

The most popular workflow scheduler to manage Hadoop jobs is arguably Apache Oozie. Like many other Hadoop products, Oozie is written in Java, and is a server-based web application that runs workflow jobs that execute Hadoop MapReduce and Pig jobs. An Oozie workflow is a collection of actions arranged in a control dependency directed acyclic graph (DAG) specified in an XML document. While Oozie has a lot of support in the Hadoop community, configuring workflows and jobs through XML attributes has a steep learning curve.

Luigi is a Python alternative, created by Spotify, that enables complex pipelines of batch jobs to be built and configured. It handles dependency resolution, workflow management, visualization, and much more. It also has a large community and supports many Hadoop technologies.

This chapter begins with the installation of Luigi and a detailed description of a workflow. Multiple examples then show how Luigi can be used to control MapReduce and Pig jobs.

Installation

Luigi is distributed through PyPI and can be installed using `pip`:

```
$ pip install luigi
```

Or it can be installed from source:

```
$ git clone https://github.com/spotify/luigi
$ python setup.py install
```

Workflows

Within Luigi, a workflow consists of a pipeline of actions, called tasks. Luigi tasks are nonspecific, that is, they can be anything that can be written in Python. The locations of input and output data for a task are known as targets. Targets typically correspond to locations of files on disk, on HDFS, or in a database. In addition to tasks and targets, Luigi utilizes parameters to customize how tasks are executed.

Tasks

Tasks are the sequences of actions that comprise a Luigi workflow. Each task declares its dependencies on targets created by other tasks. This enables Luigi to create dependency chains that ensure a task will not be executed until all of the dependent tasks and all of the dependencies for those tasks are satisfied.

Figure 5-1 depicts a workflow highlighting Luigi tasks and their dependencies.

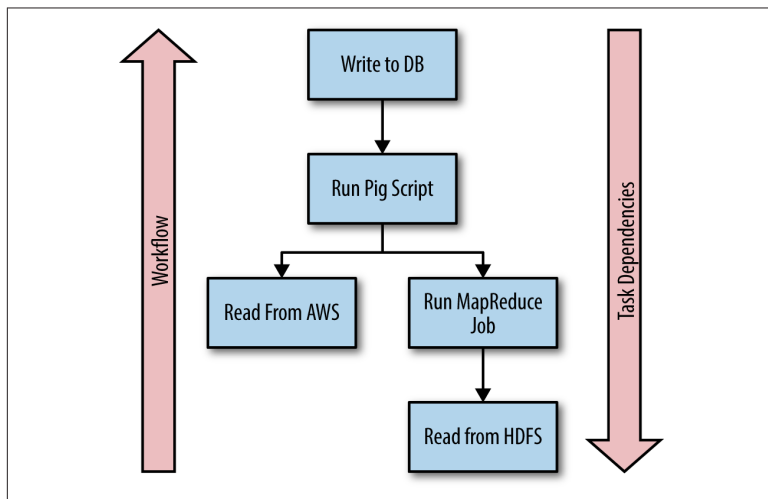


Figure 5-1. A Luigi task dependency diagram illustrates the flow of work up a pipeline and the dependencies between tasks

Target

Targets are the inputs and outputs of a task. The most common targets are files on a disk, files in HDFS, or records in a database. Luigi wraps the underlying filesystem operations to ensure that interactions with targets are atomic. This allows a workflow to be replayed from the point of failure without having to replay any of the already successfully completed tasks.

Parameters

Parameters allow the customization of tasks by enabling values to be passed into a task from the command line, programmatically, or from another task. For example, the name of a task's output may be determined by a date passed into the task through a parameter.

An Example Workflow

This section describes a workflow that implements the WordCount algorithm to explain the interaction among tasks, targets, and parameters. The complete workflow is shown in [Example 5-1](#).

Example 5-1. /python/Luigi/wordcount.py

```
import luigi

class InputFile(luigi.Task):
    """
    A task wrapping a target
    """
    input_file = luigi.Parameter()

    def output(self):
        """
        Return the target for this task
        """
        return luigi.LocalTarget(self.input_file)

class WordCount(luigi.Task):
    """
    A task that counts the number of words in a file
    """
    input_file = luigi.Parameter()
    output_file = luigi.Parameter(default='/tmp/wordcount')

    def requires(self):
```

```

The task's dependencies:
"""

return InputFile(self.input_file)

def output(self):
    """
    The task's output
    """

    return luigi.LocalTarget(self.output_file)

def run(self):
    """
    The task's logic
    """

    count = {}

    ifp = self.input().open('r')

    for line in ifp:
        for word in line.strip().split():
            count[word] = count.get(word, 0) + 1

    ofp = self.output().open('w')
    for k, v in count.items():
        ofp.write('{}\t{}\n'.format(k, v))
    ofp.close()

if __name__ == '__main__':
    luigi.run()

```

This workflow contains two tasks: `InputFile` and `WordCount`. The `InputFile` task returns the input file to the `WordCount` task. The `WordCount` tasks then counts the occurrences of each word in the input file and stores the results in the output file.

Within each task, the `requires()`, `output()`, and `run()` methods can be overridden to customize a task's behavior.

Task.requires

The `requires()` method is used to specify a task's dependencies. The `WordCount` task requires the output of the `InputFile` task:

```

def requires(self):
    return InputFile(self.input_file)

```

It is important to note that the `requires()` method cannot return a `Target` object. In this example, the `Target` object is wrapped in the `InputFile` task. Calling the `InputFile` task with the

`self.input_file` argument enables the `input_file` parameter to be passed to the `InputFile` task.

Task.output

The `output()` method returns one or more `Target` objects. The `InputFile` task returns the `Target` object that was the input for the `WordCount` task:

```
def output(self):
    return luigi.LocalTarget(self.input_file)
```

The `WordCount` task returns the `Target` object that was the output for the workflow:

```
def output(self):
    return luigi.LocalTarget(self.output_file)
```

Task.run

The `run()` method contains the code for a task. After the `requires()` method completes, the `run()` method is executed. The `run()` method for the `WordCount` task reads data from the input file, counts the number of occurrences, and writes the results to an output file:

```
def run(self):
    count = {}

    ifp = self.input().open('r')

    for line in ifp:
        for word in line.strip().split():
            count[word] = count.get(word, 0) + 1

    ofp = self.output().open('w')
    for k, v in count.items():
        ofp.write('{}\t{}\n'.format(k, v))
    ofp.close()
```

The `input()` and `output()` methods are helper methods that allow the task to read and write to `Target` objects in the `requires()` and `output()` methods, respectively.

Parameters

Parameters enable values to be passed into a task, customizing the task's execution. The WordCount task contains two parameters: `input_file` and `output_file`:

```
class WordCount(luigi.Task):
    input_file = luigi.Parameter()
    output_file = luigi.Parameter(default='/tmp/wordcount')
```

Default values can be set for parameters by using the default argument.

Luigi creates a command-line parser for each Parameter object, enabling values to be passed into the Luigi script on the command line, e.g., `--input-file input.txt` and `--output-file /tmp/output.txt`.

Execution

To enable execution from the command line, the following lines must be present in the application:

```
if __name__ == '__main__':
    luigi.run()
```

This will enable Luigi to read commands from the command line.

The following command will execute the workflow, reading from *input.txt* and storing the results in */tmp/wordcount.txt*:

```
$ python wordcount.py WordCount \
--local-scheduler \
--input-file input.txt \
--output-file /tmp/wordcount.txt
```

Hadoop Workflows

This section contains workflows that control MapReduce and Pig jobs on a Hadoop cluster.

Configuration File

The examples in this section require a Luigi configuration file, *client.cfg*, to specify the location of the Hadoop streaming jar and the path to the Pig home directory. The config files should be in the cur-

rent working directory, and an example of a config file is shown in [Example 5-2](#).

Example 5-2. python/Luigi/client.cfg

```
[hadoop]
streaming-jar: /usr/lib/hadoop-xyz/hadoop-streaming-xyz-123.jar

[pig]
home: /usr/lib/pig
```

MapReduce in Luigi

Luigi scripts can control the execution of MapReduce jobs on a Hadoop cluster by using Hadoop streaming ([Example 5-3](#)).

Example 5-3. python/Luigi/luigi_mapreduce.py

```
import luigi
import luigi.contrib.hadoop
import luigi.contrib.hdfs

class InputFile(luigi.ExternalTask):
    """
    A task wrapping the HDFS target
    """
    input_file = luigi.Parameter()

    def output(self):
        """
        Return the target on HDFS
        """
        return luigi.contrib.hdfs.HdfsTarget(self.input_file)

class WordCount(luigi.contrib.hadoop.JobTask):
    """
    A task that uses Hadoop streaming to perform WordCount
    """
    input_file = luigi.Parameter()
    output_file = luigi.Parameter()

    # Set the number of reduce tasks
    n_reduce_tasks = 1

    def requires(self):
        """
        Read from the output of the InputFile task
        """
        return InputFile(self.input_file)
```

```

def output(self):
    """
    Write the output to HDFS
    """
    return luigi.contrib.hdfs.HdfsTarget(self.output_file)

def mapper(self, line):
    """
    Read each line and produce a word and 1
    """
    for word in line.strip().split():
        yield word, 1

def reducer(self, key, values):
    """
    Read each word and produce the word and the sum of
    its values
    """
    yield key, sum(values)

if __name__ == '__main__':
    luigi.run(main_task_cls=WordCount)

```

Luigi comes packaged with support for Hadoop streaming. The task implementing the MapReduce job must subclass `luigi.contrib.hadoop.JobTask`. The `mapper()` and `reducer()` methods can be overridden to implement the map and reduce methods of a MapReduce job.

The following command will execute the workflow, reading from `/user/hduser/input.txt` and storing the results in `/user/hduser/wordcount` on HDFS:

```

$ python luigi_mapreduce.py --local-scheduler \
--input-file /user/hduser/input/input.txt \
--output-file /user/hduser/wordcount

```

Pig in Luigi

Luigi can be used to control the execution of Pig on a Hadoop cluster (Example 5-4).

Example 5-4. `python/Luigi/luigi_pig.py`

```

import luigi
import luigi.contrib.pig
import luigi.contrib.hdfs

```

```

class InputFile(luigi.ExternalTask):
    """
    A task wrapping the HDFS target
    """
    input_file = luigi.Parameter()

    def output(self):
        return luigi.contrib.hdfs.HdfsTarget(self.input_file)

class WordCount(luigi.contrib.pig.PigJobTask):
    """
    A task that uses Pig to perform WordCount
    """
    input_file = luigi.Parameter()
    output_file = luigi.Parameter()
    script_path = luigi.Parameter(default='pig/wordcount.pig')

    def requires(self):
        """
        Read from the output of the InputFile task
        """
        return InputFile(self.input_file)

    def output(self):
        """
        Write the output to HDFS
        """
        return luigi.contrib.hdfs.HdfsTarget(self.output_file)

    def pig_parameters(self):
        """
        A dictionary of parameters to pass to pig
        """
        return {'INPUT': self.input_file, 'OUTPUT': self.output_file}

    def pig_options(self):
        """
        A list of options to pass to pig
        """
        return ['-x', 'mapreduce']

    def pig_script_path(self):
        """
        The path to the pig script to run
        """
        return self.script_path

if __name__ == '__main__':
    luigi.run(main_task_cls=WordCount)

```

Luigi comes packaged with support for Pig. The task implementing the Pig job must subclass `luigi.contrib.hadoop.PigJobTask`. The `pig_script_path()` method is used to define the path to the Pig script to run. The `pig_options()` method is used to define the options to pass to the Pig script. The `pig_parameters()` method is used to pass parameters to the Pig script.

The following command will execute the workflow, reading from `/user/hduser/input.txt` and storing the results in `/user/hduser/output` on HDFS. The `--script-path` parameter is used to define the Pig script to execute:

```
$ python luigi_pig.py --local-scheduler \  
--input-file /user/hduser/input/input.txt \  
--output-file /user/hduser/output \  
--script-path pig/wordcount.pig
```

Chapter Summary

This chapter introduced Luigi as a Python workflow scheduler. It described the components of a Luigi workflow and contained examples of using Luigi to control MapReduce jobs and Pig scripts.

About the Authors

Zachary Radtka is a platform engineer at the data science firm Miner & Kasch and has extensive experience creating custom analytics that run on petabyte-scale datasets. Zach is an experienced educator, having instructed collegiate-level computer science classes, professional training classes on Big Data technologies, and public technology tutorials. He has also created production-level analytics for many industries, including US government, financial, healthcare, telecommunications, and retail.

Donald Miner is founder of the data science firm Miner & Kasch, and specializes in Hadoop enterprise architecture and applying machine learning to real-world business problems.

Donald is the author of the O'Reilly book *MapReduce Design Patterns* and the upcoming O'Reilly book *Enterprise Hadoop*. He has architected and implemented dozens of mission-critical and large-scale Hadoop systems within the US government and Fortune 500 companies. He has applied machine learning techniques to analyze data across several verticals, including financial, retail, telecommunications, health care, government intelligence, and entertainment. His PhD is from the University of Maryland, Baltimore County, where he focused on machine learning and multiagent systems. He lives in Maryland with his wife and two young sons.
